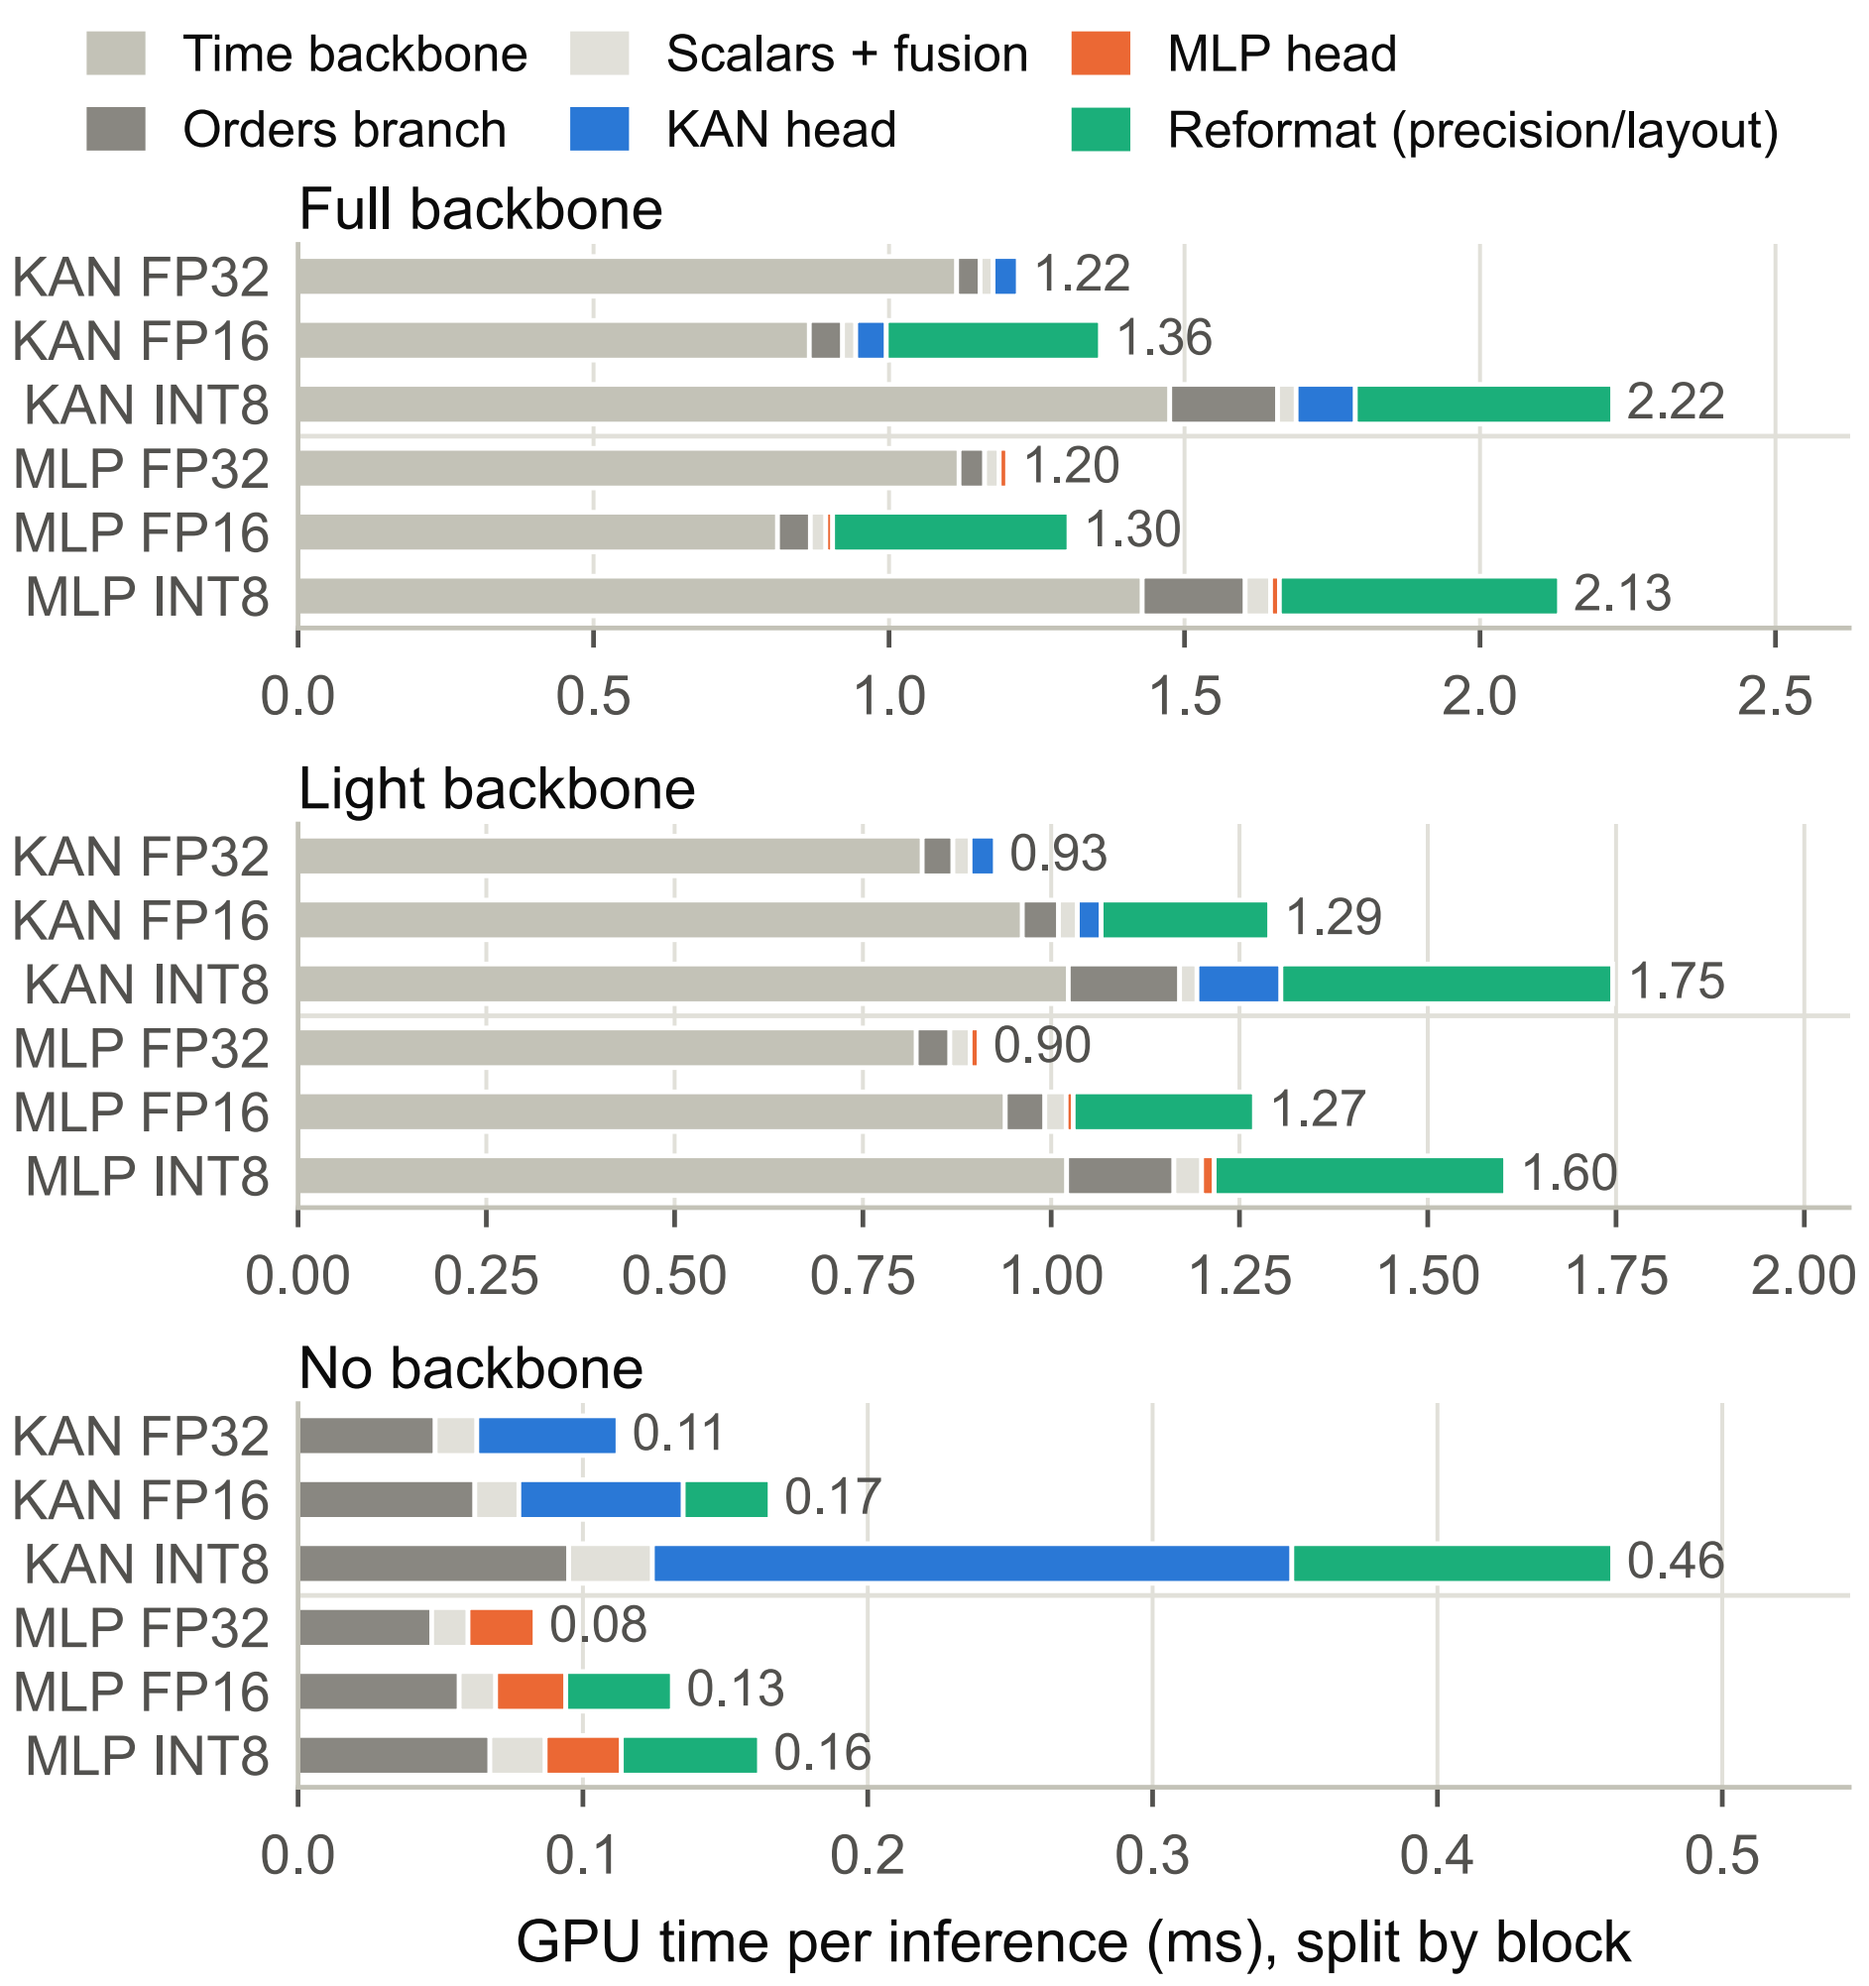
Time backbone
Scalars + fusion
MLP head
Orders branch
KAN head
Reformat (precision/layout)
Full backbone
KAN FP32
KAN FP16
KAN INT8
MLP FP32
MLP FP16
MLP INT8
1.22
1.36
2.22
1.20
1.30
2.13
0.0
0.5
1.0
1.5
2.0
2.5
Light backbone
KAN FP32
KAN FP16
KAN INT8
MLP FP32
MLP FP16
MLP INT8
0.93
1.29
1.75
0.90
1.27
1.60
0.00
0.25
0.50
0.75
1.00
1.25
1.50
1.75
2.00
No backbone
KAN FP32
KAN FP16
KAN INT8
MLP FP32
MLP FP16
MLP INT8
0.11
0.17
0.46
0.08
0.13
0.16
0.0
0.1
0.2
0.3
0.4
0.5
GPU time per inference (ms), split by block